


Convolutional Neural Networks (CNNs)

- handles image data well

 - ↳ image classification CIFAR-10

 - ↳ object detection

 - ↳ neural style transfer

- pass filters over pixels

- CNNs work well when structure exists in our data

 - ↳ spatial, temporal, semantic

CNN's Structure

goal:
reduce
of params
needed to
learn
in FC
layers.

1) Convolutional

↳ # of filters

↳ as # of filters ↑, # of params ↑

↳ size of filters

↳ as size of filters increases, # of params ↑

↳ padding / stride
level of detail ↓

2) Pooling

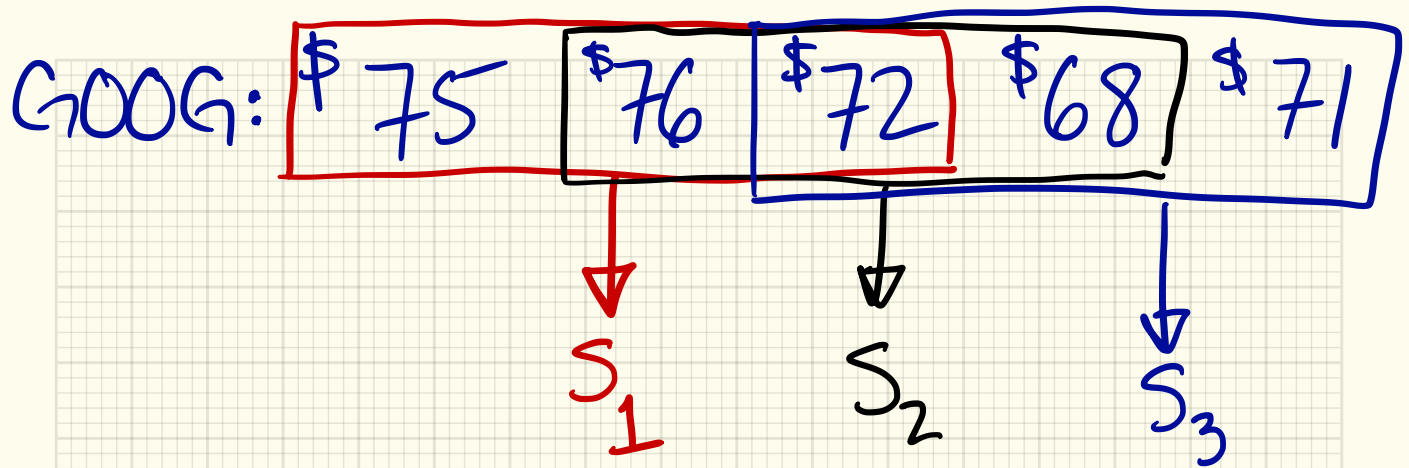
Flatten(·)

3) Fully / Densely Connected

↳ (FC)

w_1	w_2	w_3
w_4	w_5	w_6
w_7	w_8	w_9

"Parameter Sharing"




1D convolution / filter of length 3.

$$\text{filter} = [w_1, w_2, w_3]$$

tweet: "We have a Commander in Cheese
-- I mean, Chief."

t_1 t_2 t_3 ...

1D Convolution w/ filter of size 2.



3D Convolution $\rightarrow 3 \times 3 \times 5$

filter $\Rightarrow \underline{a} \times \underline{b} \times \underline{c}$